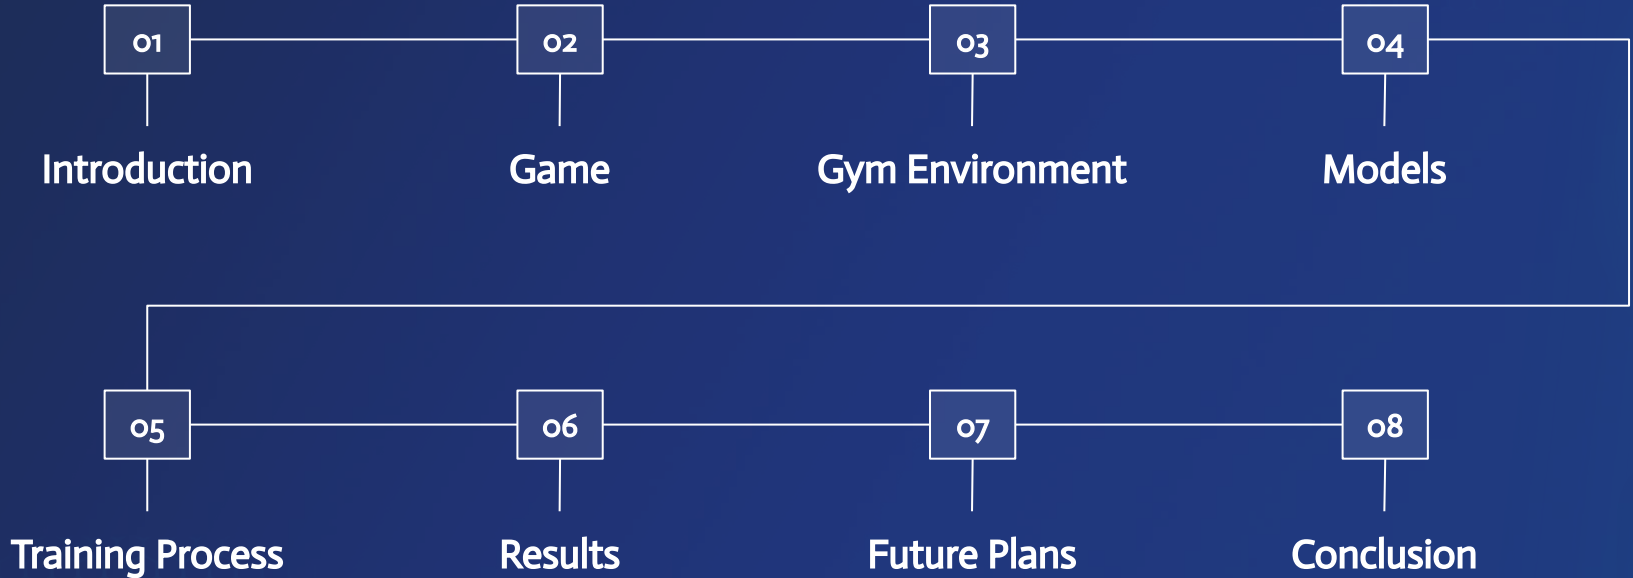


# Air-Hockey Bot



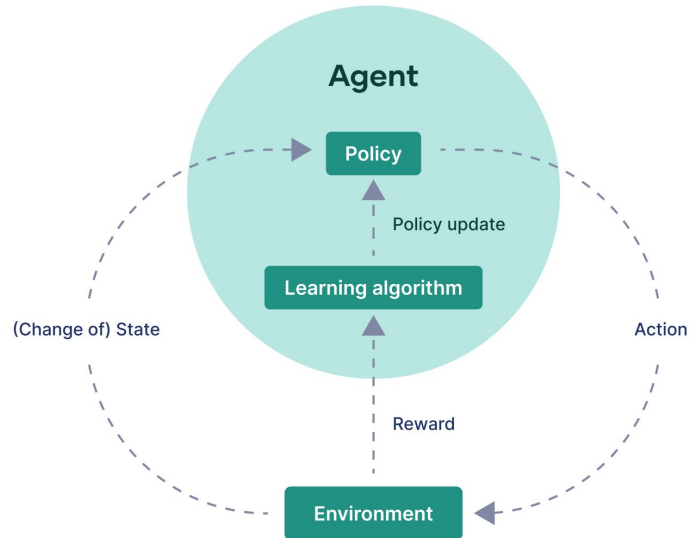
Jorge Bris Moreno, Eric Dapkus, Brian Kwon, Kang Liu, and Billy McGloin  
DSAN 6600, Final Project, April 24, 2024

# Contents




# Introduction

## The general framework of reinforcement learning



# Air Hockey Game



# Environment

[Deck.html](#)

(<https://kang.georgetown.domains/dsan-website/66oo-website/project-presentation/deck.html#/title-slide>)



# Models

# Convolutional Neural Network

```
class DQN(nn.Module):
    def __init__(self, outputs):
        super(DQN, self).__init__()
        self.conv1 = nn.Conv2d(3, 16, kernel_size=3, stride=1, padding=1) # Small kernel, stride 1
        self.bn1 = nn.BatchNorm2d(16)
        self.pool1 = nn.AvgPool2d(2) # Pooling to reduce dimension

        self.conv2 = nn.Conv2d(16, 32, kernel_size=3, stride=1, padding=1) # Small kernel, stride 1
        self.bn2 = nn.BatchNorm2d(32)
        self.pool2 = nn.AvgPool2d(2) # Pooling to reduce dimension

        # Calculate the size of the output from the pooling layer
        self._to_linear = None
        self._get_conv_output([1, 3, 16, 8]) # Input shape sample

        self.head = nn.Linear(self._to_linear, outputs) # Linear layer to output the Q-values for each action

    def _get_conv_output(self, shape):
        input = torch.rand(shape)
        output = self.pool1(self.bn1(self.conv1(input)))
        output = self.pool2(self.bn2(self.conv2(output)))
        self._to_linear = int(torch.numel(output) / output.shape[0])

    def forward(self, x):
        x = F.relu(self.pool1(self.bn1(self.conv1(x))))
        x = F.relu(self.pool2(self.bn2(self.conv2(x))))
        x = x.view(x.size(0), -1) # Flatten the features for the linear layer
        return self.head(x)
```

# Linear Neural Network

```
# Linear Neural network model
class DQN_linear(nn.Module):
    def __init__(self, input_dim, output_dim):
        super(DQN_linear, self).__init__()
        self.linear1 = nn.Linear(input_dim, 16) # First hidden layer
        self.linear2 = nn.Linear(16, 32) # Second hidden layer
        self.linear3 = nn.Linear(32, 16) # Third hidden layer
        self.head = nn.Linear(16, output_dim) # Output layer to produce the final outputs

    def forward(self, x):
        x = F.relu(self.linear1(x)) # Activation function for first layer
        x = F.relu(self.linear2(x)) # Activation function for second layer
        x = F.relu(self.linear3(x)) # Activation function for third layer
        x = self.head(x) # Output layer does not need an activation for Q-value estimation
        return x
```



# CNN & Linear Combo

```
class DQN(nn.Module):
    def __init__(self, outputs):
        super(DQN, self).__init__()
        self.conv1 = nn.Conv2d(3, 16, kernel_size=4) # Small kernel, stride 1
        self.bn1 = nn.BatchNorm2d(16)

        self.conv2 = nn.Conv2d(16, 32, kernel_size=4) # Small kernel, stride 1
        self.bn2 = nn.BatchNorm2d(32)

        # Calculate the size of the output
        self._to_linear = None
        self._get_conv_output([1, 3, 16, 8]) # Input shape sample

        self.linear = nn.Linear(self._to_linear, 18) # Linear layer to output the Q-values for each action
        self.head = nn.Linear(18, outputs) # Head layer to output the final Q-values

    def _get_conv_output(self, shape):
        input = torch.rand(shape)
        output = self.bn1(self.conv1(input))
        output = self.bn2(self.conv2(output))
        self._to_linear = int(torch.numel(output) / output.shape[0])

    def forward(self, x):
        x = F.relu(self.bn1(self.conv1(x)))
        x = F.relu(self.bn2(self.conv2(x)))
        x = x.view(x.size(0), -1) # Flatten the features before the linear layer
        x = F.relu(self.linear(x))
        return self.head(x)
```

# Training Process (Training environment)

## Google Colab Pro

- L4 GPU
- 12.7GB system ram
- 4.1 credit per hour



```
#check GPU
gpu_info = !nvidia-smi
gpu_info = '\n'.join(gpu_info)
if gpu_info.find('failed') >= 0:
    print('Not connected to a GPU')
else:
    print(gpu_info)
```

Google Colab

```
#install required packages for colab
!pip install gymnasium
#run code
!python dqn.py
```

## Intel NUC


- 4 x NUC11
- i5
- 64GB ram
- 500 GB m.2



## Ubuntu Server (22.04.4 LTS)

### Software

- Screen
- miniconda
- pip



Hostnames: James, Ben, Jeff, Nakul

# Training Process Cont. (Training Data)

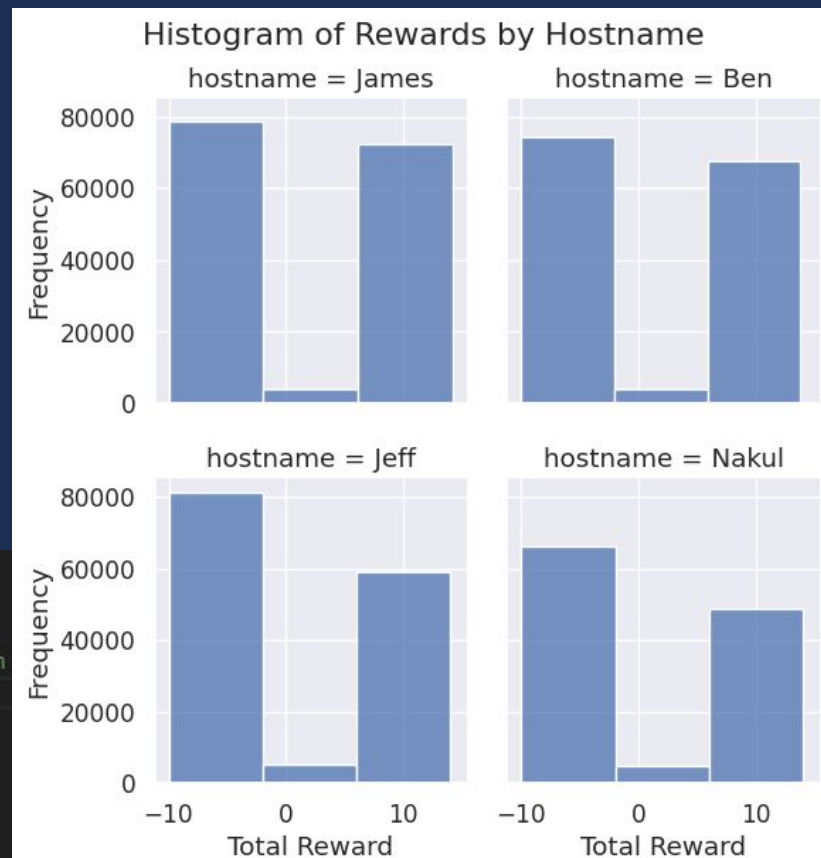
james (155257 episodes)  
- Opponent new bot logic

ben (146760 episodes)  
- Opponent new bot logic  
- TARGET\_UPDATE = 750

jeff (145268 episodes)  
- Opponent old bot logic

nakul  
- Opponent old bot logic  
- TARGET\_UPDATE = 750

```
# hyperparameters
BATCH_SIZE = 64
GAMMA = 0.99 # Discount factor for future rewards
EPS_START = 0.95 # Initial epsilon value for epsilon-greedy action selection
EPS_END = 0.05 # Final epsilon value for epsilon-greedy action selection
EPS_DECAY = 80000 # Rate at which epsilon decreases
LR = 1e-4 # Learning rate for the optimizer
TARGET_UPDATE = 1000 # How often to update the target network
MEMORY_CAPACITY = 10000 # Capacity of the replay memory
NUM_EPISODES = 500000 # Number of episodes to train
```





# Training Process Cont. (Training Data)





# Proof of Concept











# Air Hockey Implementation





# Future Plans

1. Extend Training Duration
2. Evaluate Alternate Rewards
3. Human v. RL?
4. Implement Supervised Learning
5. RL v. RL?
6. Upgrade RL Network



# Conclusion



# Thanks!

Special thanks to our professors!

Reference:

Google Colab. <https://colab.research.google.com/#>

**CREDITS:** This presentation template was created by [Slidesgo](#), and includes icons by [Flaticon](#), and infographics & images by [Freepik](#)

# Questions?

